

Homework2 – Introduction to Artificial Neural Networks with Keras

Building an Image Classifier

First let's import TensorFlow and Keras.

```
In [1]: ▶ import tensorflow as tf
        from tensorflow import keras
```

```
In [2]: ▶ import os
        import gzip
        import numpy as np

        current_dir = os.getcwd() # First, get the path of the working directory
        path = current_dir+'\\Fashion_MNIST_Data'

        # Second, import the Fashion MNIST data from the current directory+'\\Fashion
        f = gzip.open(path+'\\train-labels-idx1-ubyte.gz','rb') #Load the training l
        y_train_full=np.frombuffer(f.read(), dtype=np.uint8,offset=8) # due to header
        f.close()

        f = gzip.open(path+'\\train-images-idx3-ubyte.gz','rb') #Load the training a
        X_train_full=np.frombuffer(f.read(), dtype=np.uint8,offset=16) # due to heade
        f.close()
        X_train_full=X_train_full.reshape(len(y_train_full), 784)
        X_train_full1=X_train_full.reshape(len(y_train_full), 28,28)

        f = gzip.open(path+'\\t10k-labels-idx1-ubyte.gz','rb') #Load the test label
        y_test=np.frombuffer(f.read(), dtype=np.uint8,offset=8)
        f.close()

        f = gzip.open(path+'\\t10k-images-idx3-ubyte.gz','rb') #Load the test data
        X_test=np.frombuffer(f.read(), dtype=np.uint8,offset=16).reshape(len(y_test),
        f.close()
```

```
In [3]: ▶ # Check the consistency
print(len(X_train_full))
print(len(y_train_full))
print(len(X_test))
print(len(y_test))
y_test
```

```
60000
60000
10000
10000
```

```
Out[3]: array([9, 2, 1, ..., 8, 1, 5], dtype=uint8)
```

The training set contains 60,000 grayscale images, each 28×28 pixels. The class labels are:

Label: Description

- 0: TT-shirt/top
- 1: Trouser
- 2: Pullover
- 3: Dress
- 4: Coat
- 5: Sandal
- 6: Shirt
- 7: Sneaker
- 8: Bag
- 9: Ankle boot

```
In [4]: ▶ #Check the shape of the dataset
print(X_train_full.shape)
print(y_train_full.shape)
print(X_test.shape)
```

```
(60000, 784)
(60000,)
(10000, 784)
```

Each pixel intensity is represented as a byte (0 to 255):

Let's split the full training set into a validation set and a (smaller) training set. We also scale the pixel intensities down to the 0-1 range and convert them to floats, by dividing by 255. .

```
In [6]: ▶ X_valid, X_train = X_train_full[:5000] / 255., X_train_full[5000:] / 255.
y_valid, y_train = y_train_full[:5000], y_train_full[5000:]
X_test = X_test / 255.
```

Here are the corresponding class names:

```
In [7]: class_names = ["T-shirt/top", "Trouser", "Pullover", "Dress", "Coat", "Sandal",
                        "Shirt", "Sneaker", "Bag", "Ankle boot"]
```

So the first image in the training set is a coat:

```
In [8]: class_names[y_train[0]]
```

```
Out[8]: 'Coat'
```

Let's take a look at a sample of the images in the dataset:

```
In [11]: import matplotlib.pyplot as plt
n_rows = 5
n_cols = 8
plt.figure(figsize=(n_cols * 1.2, n_rows * 1.2)) # scale up, otherwise, 28x28
for row in range(n_rows):
    for col in range(n_cols):
        index = n_cols * row + col
        plt.subplot(n_rows, n_cols, index + 1)
        plt.imshow(X_train_full[index], cmap="binary", interpolation="nearest")
        #plt.axis('off')
        plt.title(class_names[y_train_full[index]], fontsize=12)
plt.subplots_adjust(wspace=0.2, hspace=1.0)

plt.show()
```



Figure 2: Samples from Fashion MNIST in a 4 x 10 grid.

Model 1





Fig. 6: Handling multiple inputs.

```
In [12]: ▶ np.random.seed(42)
          tf.random.set_seed(42)
```


```
In [13]: ▶ input_A = keras.layers.Input(shape=[400], name="wide_input")
          input_B = keras.layers.Input(shape=[400], name="deep_input")
          hidden1 = keras.layers.Dense(30, activation="relu")(input_B)
          hidden2 = keras.layers.Dense(30, activation="relu")(hidden1)
          concat = keras.layers.concatenate([input_A, hidden2])
          output = keras.layers.Dense(10, name="output", activation="softmax")(concat)
          model1 = keras.models.Model(inputs=[input_A, input_B], outputs=[output])
```

```
In [14]: ▶ model1.compile(loss=keras.losses.sparse_categorical_crossentropy,
                        optimizer=keras.optimizers.SGD(),
                        metrics=[keras.metrics.sparse_categorical_accuracy])
          checkpoint_cb = keras.callbacks.ModelCheckpoint("my_keras_model1.h5", save_be
```

In [15]:  model1.summary()

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
=====			
deep_input (InputLayer)	[(None, 400)]	0	[]
dense (Dense)	(None, 30)	12030	['deep_in ut[0][0]']
wide_input (InputLayer)	[(None, 400)]	0	[]
dense_1 (Dense)	(None, 30)	930	['dense[0] [0]']
concatenate (Concatenate)	(None, 430)	0	['wide_in ut[0][0]', 'dense_1 [0][0]']
output (Dense)	(None, 10)	4310	['concaten ate[0][0]']
=====			
Total params: 17,270			
Trainable params: 17,270			
Non-trainable params: 0			

In [16]: 

```

X_train_A, X_train_B = X_train[:, :400], X_train[:, 384:]
X_valid_A, X_valid_B= X_valid[:, :400], X_valid[:, 384:]
X_test_A, X_test_B = X_test[:, :400], X_test[:, 384:]

```

```
In [17]: history1 = model1.fit((X_train_A, X_train_B), y_train, epochs=2000,
                             validation_data=((X_valid_A, X_valid_B), y_valid), callback=callbacks)

sparse_categorical_accuracy: 0.8222 - val_loss: 0.4898 - val_sparse_categorical_accuracy: 0.8336
Epoch 5/2000
1719/1719 [=====] - 6s 3ms/step - loss: 0.4893 - sparse_categorical_accuracy: 0.8294 - val_loss: 0.4658 - val_sparse_categorical_accuracy: 0.8424
Epoch 6/2000
1719/1719 [=====] - 5s 3ms/step - loss: 0.4712 - sparse_categorical_accuracy: 0.8350 - val_loss: 0.4524 - val_sparse_categorical_accuracy: 0.8470
Epoch 7/2000
1719/1719 [=====] - 5s 3ms/step - loss: 0.4573 - sparse_categorical_accuracy: 0.8407 - val_loss: 0.4413 - val_sparse_categorical_accuracy: 0.8502
Epoch 8/2000
1719/1719 [=====] - 6s 3ms/step - loss: 0.4454 - sparse_categorical_accuracy: 0.8449 - val_loss: 0.4448 - val_sparse_categorical_accuracy: 0.8428
Epoch 9/2000
1719/1719 [=====] - 5s 3ms/step - loss: 0.4353 -
```


Loading saved best version of the models and evaluating them using the test dataset.

```
In [21]: model1 = keras.models.load_model("my_keras_model1.h5") # rollback to best model

In [22]: y_pred_main= model1.evaluate((X_test_A, X_test_B),y_test)

313/313 [=====] - 1s 2ms/step - loss: 0.3758 - sparse_categorical_accuracy: 0.8682
```

```
In [23]: import pandas as pd
X=pd.DataFrame(history1.history)
X=X[:200]
X.plot(figsize=(8, 5))
plt.grid(True)
plt.gca().set_ylim(0, 1)
plt.show()
```



Model 2





Fig. 7: Handling multiple inputs.

```
In [12]: ▶ np.random.seed(42)
         tf.random.set_seed(42)
```


```
In [13]: ▶ input_A = keras.layers.Input(shape=[400], name="wide_input")
         hidden1 = keras.layers.Dense(30, activation="relu")(input_A)
         input_B = keras.layers.Input(shape=[400], name="deep_input")
         hidden11 = keras.layers.Dense(30, activation="relu")(input_B)
         hidden12 = keras.layers.Dense(30, activation="relu")(hidden11)
         concat = keras.layers.concatenate([hidden1, hidden12])
         output = keras.layers.Dense(10, name="output", activation="softmax")(concat)
         model2 = keras.models.Model(inputs=[input_A, input_B], outputs=[output])
```

```
In [14]: ▶ model2.compile(loss=keras.losses.sparse_categorical_crossentropy,
                        optimizer=keras.optimizers.SGD(),
                        metrics=[keras.metrics.sparse_categorical_accuracy])
         checkpoint_cb = keras.callbacks.ModelCheckpoint("my_keras_model1.h5", save_be
```


In [15]:  model2.summary()

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
=====			
deep_input (InputLayer)	[(None, 400)]	0	[]
wide_input (InputLayer)	[(None, 400)]	0	[]
dense_1 (Dense)	(None, 30)	12030	['deep_inp ut[0][0]']
dense (Dense)	(None, 30)	12030	['wide_inp ut[0][0]']
dense_2 (Dense)	(None, 30)	930	['dense_1 [0][0]']
concatenate (Concatenate)	(None, 60)	0	['dense[0] [0]', 'dense_2 [0][0]']
output (Dense)	(None, 10)	610	['concaten ate[0][0]']
=====			
Total params: 25,600			
Trainable params: 25,600			
Non-trainable params: 0			

In [16]: 

```

X_train_A, X_train_B = X_train[:, :400], X_train[:, 384:]
X_valid_A, X_valid_B = X_valid[:, :400], X_valid[:, 384:]
X_test_A, X_test_B = X_test[:, :400], X_test[:, 384:]
X_new_A, X_new_B = X_test_A[:1000], X_test_B[:1000]

```

```
In [17]: history2 = model2.fit((X_train_A, X_train_B), y_train, epochs=2000,
                             validation_data=((X_valid_A, X_valid_B), y_valid), callback=callbacks)
```

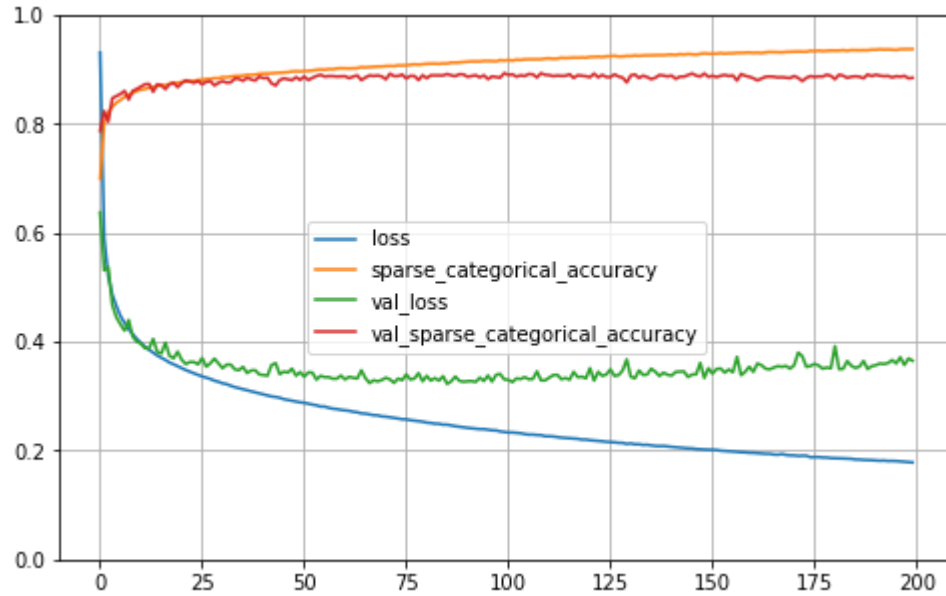
```
Epoch 1560/2000
1719/1719 [=====] - 6s 4ms/step - loss: 0.0246 - sparse_categorical_accuracy: 0.9929 - val_loss: 1.3215 - val_sparse_categorical_accuracy: 0.8562
Epoch 1561/2000
1719/1719 [=====] - 6s 4ms/step - loss: 0.0239 - sparse_categorical_accuracy: 0.9933 - val_loss: 1.3333 - val_sparse_categorical_accuracy: 0.8590
Epoch 1562/2000
1719/1719 [=====] - 6s 4ms/step - loss: 0.0229 - sparse_categorical_accuracy: 0.9933 - val_loss: 1.3263 - val_sparse_categorical_accuracy: 0.8574
Epoch 1563/2000
1719/1719 [=====] - 6s 4ms/step - loss: 0.0225 - sparse_categorical_accuracy: 0.9936 - val_loss: 1.3143 - val_sparse_categorical_accuracy: 0.8622
Epoch 1564/2000
1719/1719 [=====] - 6s 4ms/step - loss: 0.0264 - sparse_categorical_accuracy: 0.9921 - val_loss: 1.3089 - val_sparse_categorical_accuracy: 0.8632
```

Loading saved best version of the models and evaluating them using the test dataset.

```
In [18]: model2 = keras.models.load_model("my_keras_model2.h5") # rollback to best model
y_pred_main2 = model2.evaluate((X_test_A, X_test_B), y_test)
```

```
313/313 [=====] - 1s 2ms/step - loss: 1.6682 - sparse_categorical_accuracy: 0.8536
```

```
In [19]: import pandas as pd
X2=pd.DataFrame(history2.history)
X2=X2[:200]
X2.plot(figsize=(8, 5))
plt.grid(True)
plt.gca().set_ylim(0, 1)
plt.show()
```



Model 3




Fig. 7: Handling multiple inputs.

```
In [12]: np.random.seed(42)
tf.random.set_seed(42)
```

```
In [13]: ▶ input_A = keras.layers.Input(shape=[300], name="wide_input")
input_B = keras.layers.Input(shape=[300], name="deep_input1")
input_c = keras.layers.Input(shape=[300], name="deep_input2")
hidden1 = keras.layers.Dense(30, activation="relu")(input_A)
hidden11 = keras.layers.Dense(30, activation="relu")(input_B)
hidden12 = keras.layers.Dense(30, activation="relu")(hidden11)
hidden21 = keras.layers.Dense(30, activation="relu")(input_c)
hidden22 = keras.layers.Dense(30, activation="relu")(hidden21)
hidden33 = keras.layers.Dense(30, activation="relu")(hidden22)

concat = keras.layers.concatenate([hidden1, hidden12,hidden33])
output = keras.layers.Dense(10, name="output",activation="softmax")(concat)
model3 = keras.models.Model(inputs=[input_A, input_B,input_c], outputs=[output])
```

```
In [14]: ▶ X_train_A, X_train_B ,X_train_C = X_train[:, :300], X_train[:, 243:543],X_train[:, 543:]
X_valid_A, X_valid_B ,X_valid_C = X_valid[:, :300], X_valid[:, 243:543],X_valid[:, 543:]
X_test_A, X_test_B,X_test_C= X_test[:, :300], X_test[:, 243:543],X_test[:, 543:]

X_new_A, X_new_B, X_new_C= X_test_A[:1000], X_test_B[:1000],X_test_C[:1000]
```

```
In [28]: ▶ model3.compile(loss=keras.losses.sparse_categorical_crossentropy,
                        optimizer=keras.optimizers.SGD(),
                        metrics=[keras.metrics.sparse_categorical_accuracy])
checkpoint_cb = keras.callbacks.ModelCheckpoint("my_keras_model3.h5", save_best_only=True)
```

In [29]: `model3.summary()`

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
=====			
deep_input2 (InputLayer)	[(None, 300)]	0	[]
deep_input1 (InputLayer)	[(None, 300)]	0	[]
dense_3 (Dense)	(None, 30)	9030	['deep_inp ut2[0][0]']
wide_input (InputLayer)	[(None, 300)]	0	[]
dense_1 (Dense)	(None, 30)	9030	['deep_inp ut1[0][0]']
dense_4 (Dense)	(None, 30)	930	['dense_3 [0][0]']
dense (Dense)	(None, 30)	9030	['wide_inp ut[0][0]']
dense_2 (Dense)	(None, 30)	930	['dense_1 [0][0]']
dense_5 (Dense)	(None, 30)	930	['dense_4 [0][0]']
concatenate (Concatenate)	(None, 90)	0	['dense[0] [0]', 'dense_2 [0][0]', 'dense_5 [0][0]']
output (Dense)	(None, 10)	910	['concaten ate[0][0]']

```

=====
Total params: 30,790
Trainable params: 30,790
Non-trainable params: 0

```

```
In [31]: history3 = model3.fit((X_train_A, X_train_B,X_train_C), y_train, epochs=1700,
                             validation_data=((X_valid_A, X_valid_B,X_valid_C), y_vali
```

```
sparse_categorical_accuracy: 1.0000 - val_loss: 1.4241 - val_sparse_categ
orical_accuracy: 0.8770
Epoch 1251/1700
1719/1719 [=====] - 7s 4ms/step - loss: 0.0019 -
sparse_categorical_accuracy: 1.0000 - val_loss: 1.4267 - val_sparse_categ
orical_accuracy: 0.8772
Epoch 1252/1700
1719/1719 [=====] - 7s 4ms/step - loss: 0.0016 -
sparse_categorical_accuracy: 1.0000 - val_loss: 1.4599 - val_sparse_categ
orical_accuracy: 0.8762
Epoch 1291/1700
1719/1719 [=====] - 7s 4ms/step - loss: 0.0017 -
sparse_categorical_accuracy: 1.0000 - val_loss: 1.4558 - val_sparse_categ
orical_accuracy: 0.8768
Epoch 1292/1700
1719/1719 [=====] - 7s 4ms/step - loss: 0.0016 -
sparse_categorical_accuracy: 1.0000 - val_loss: 1.4566 - val_sparse_categ
orical_accuracy: 0.8762
Epoch 1293/1700
```

```
In [ ]:
```

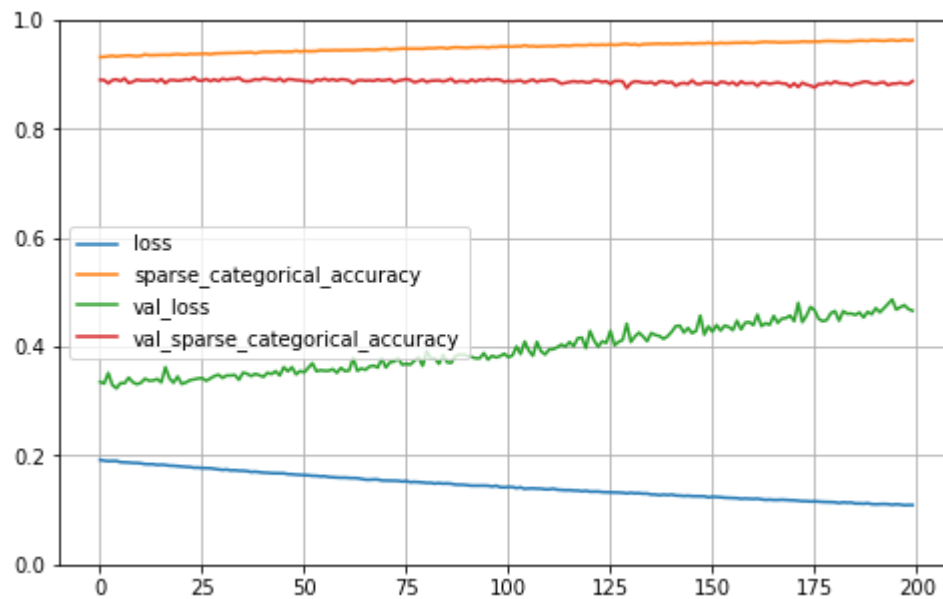
NOTE; my third model was always crushig , hence i had to reduce the number of epoch to 1700.

Loading saved best version of the models and evaluating them using the test dataset.

```
In [32]: model3 = keras.models.load_model("my_keras_model3.h5") # rollback to best mod
y_pred_main3= model3.evaluate((X_test_A, X_test_B,X_test_C),y_test)
```

```
313/313 [=====] - 3s 3ms/step - loss: 0.3389 - spa
rse_categorical_accuracy: 0.8814
```

```
In [34]: import pandas as pd
X3=pd.DataFrame(history3.history)
X3=X3[:200]
X3.plot(figsize=(8, 5))
plt.grid(True)
plt.gca().set_ylim(0, 1)
plt.show()
```



```
In [ ]:
```